


DOMAIN 04 · SIX SCENARIOS · RE-RUN AT EVERY ROADMAP MILESTONE AND EVERY PARTNER PHASE GATE

Execution and Platform Risk

Does one engine, nine layers and rules held as writing survive contact with real feeds, a second rule maker and a delivery date?

Nine layers, six places where two parts touch, seven things in the evidence store and five screens. Two of the nine layers are built by senior hands inside the company and are never handed to a build partner: the store that is only added to, and the two dates that say which day a record is about and which day it was written. **The whole cost structure rests on one sentence, which is that everything above the platform is content, a setting or a list of fields.** A platform claim is proven at the moment it is tested rather than the moment it is drawn, and the test is scheduled for Q3 2028, on the milestone the seed is priced against.



Bubble size is the risk score, likelihood 1 to 3 multiplied by impact 1 to 3. The number inside is the scenario row in the table overleaf.

High 7 to 9

Medium 4 to 6

Low 1 to 3

SOURCES ON THIS PAGE

The architecture	4orm Finance Technical Architecture 03.2, 17 August 2026: nine layers and who builds each, six boundaries and what must never cross them, the seven things in the evidence store, the record shape, and the four things that have to be right in the first release. Read on a countersigned mutual non-disclosure agreement.	MODEL
The open tier	4orm Finance Technology Overview 03.1, 17 August 2026: what the software does, what it refuses to do, and how a reader can test the refusal.	MODEL
Scoring	Likelihood and impact are management judgements on a 1 to 3 scale.	MODEL

The six scenarios, scored

#	Scenario and trigger	L	I	Score	Early warning signal
1	One of the four things that cannot be added later is traded for a date The store that is only added to, the two dates on every record, rules held as writing, and customers kept apart at the key. All four cost days or weeks in the first release and are close to out of reach afterwards, and all four are invisible in a demonstration	3	3	9	A release plan moves any of the four behind a screen, or one administrative change to a stored value is granted under pressure on a support call
2	The second rule maker needs the data structure rebuilt Property management trust reconciliation is scheduled to go live on the same engine in Q3 2028. If it needs a second version of the software, this is several products wearing one name	2	3	6	Any line item in the second rule maker estimate reads as a new field, a new workflow state or a new binding rule rather than as content
3	The held mechanism delays the two layers nobody else may build The record shape names nine fields. A tenth and eleventh, where the chain is anchored, how a replacing record is admitted and how two exports a year apart land on identical bytes are all held pending a filing with counsel	2	3	6	The architecture note is due and counsel has not confirmed the filing, or the build partner asks for the held items in order to proceed
4	Re-export does not reproduce the same bytes The whole claim that a record has gone unaltered rests on exporting a period again a year later and getting identical bytes. Milestone 10 puts an automated test on every build rather than in a questionnaire	2	2	4	The re-export test is described rather than running, or it runs against synthetic fixtures rather than production shaped data
5	The assistant writes into the record It may search inside one customer with every answer tied to a record identifier. It may never put a made up value into the record. It finds and drafts, and a person signs	2	2	4	Any log shows an assistant draft reaching a filing or an examiner facing screen without a logged human approval in front of it
6	Ingest stops being configuration Bank, custodian and trust feeds are built by the partner and reviewed. It is the layer most likely of the nine to accumulate work written for one customer, and connectors are metered as revenue inside the least evidenced line	2	2	4	Average engineering hours for each new feed rises rather than falls across 2027, or a field that will not read is defaulted instead of failing the record

Legend. Likelihood and impact each run 1 to 3 and the score is their product. Five of the six can be tested by asking to watch something run rather than by reading a description of it, which is the only useful kind of assurance at this stage.

The design is set by what it refuses, and a refusal is the only part of an architecture a reviewer can actually test.

Six places two parts touch, and nowhere else. A reviewer can test the whole claim by asking for the seventh place. If one exists and is not on the drawing, the drawing is wrong and the design is weaker than it reads. That is a fair question to put in the first call, and it is the shape every scenario below is scored against.

1

One of the four things that cannot be added later is traded for a date

SCORE 9 · L3 × I3

EVIDENCE · TWENTY THREE PRODUCTS, AND WHAT NONE OF THEM CAN DO

The assessed set is the argument for building the invisible layers first. Zero of twenty three carry the Canadian duty, zero prove a record has gone unaltered and zero reconstruct history they were not connected for: all twenty three start at the connection date. Two vendors use the word immutable. One describes no mechanism at all and the other means bookkeeping that is only added to. Four compute a daily position on money held for customers and all four are British. **The market has already made this trade, repeatedly, and that is precisely why the category is open.** Every one of those products shipped screens that demonstrate well. None of them shipped the part an examiner opens with.

RESPONSE FOR 4ORM

The architecture prices all four honestly. The store that is only added to costs days now. The two dates cost one to two weeks. Rules held as writing cost two to three weeks. Customers kept apart at the key cost days if it is designed in. Added later, the first leaves the oldest and most examined records with nothing behind them, the second means the rule versions that applied to past decisions were never kept and cannot be worked out from anything, the third means rules already in code have to be pulled out of it and every customer moved, and the fourth means re-keying live customer data, which is the most likely place for a breach to be let in and it would be let in by the fix. **The three ways they get dropped are named, in order of likelihood: a demonstration date, a support call, and a sensible engineer.** The controls are equally concrete. Put all four in the architecture note that gates the third instalment. Refuse the administrative change on the support call, once, in writing, before the first one is asked for. And require each of the four to be demonstrated running rather than described: the same bytes on a second export, any day run again giving the same answer, a duty shipped with no code release, and a cross customer read returning nothing.

SOURCES ON THIS PAGE

The four things	Cost now, what adding each later means, and the test for each. 4orm Finance Technical Architecture 03.2, 17 August 2026, the four things that have to be right in the first release, and the three ways they get dropped.	MODEL
The assessed set	0 of 23 carry the Canadian duty, 0 prove a record unaltered, 0 reconstruct history and all 23 start at the connection date. Two use the word immutable, neither describes a mechanism. 4orm Finance Comparables and Benchmarks, 17 August 2026 · 4orm Finance Service and Price Comparison, 17 August 2026	MODEL
The four that compute a daily position, all British	AutoRek · Kani Payments · Guardexia · Equali . Read on 16 and 17 August 2026.	PRIMARY
What a reviewer must be able to trace	End user ledger, reconciliation, reliance boundaries and the evidence a reviewer must be able to open and trace. Bank of Canada safeguarding guideline	PRIMARY
The gating milestone	Architecture note delivered and accepted, Q3 2026, tested against the four properties in the architecture pack, before the phase one deposit. 4orm Finance master pro forma rebuild of 16 August 2026, base case, recalculated in full, milestones tab.	MODEL

2 The second rule maker needs the data structure rebuilt

SCORE 6 · L2 × I3

EVIDENCE · WHAT SITS ON THE OTHER SIDE OF THE Q3 2028 TEST

The second rule maker is not a slide. It carries \$698,381 in 2028 rising to \$13,766,769 by 2031 and \$26,049,518 across the plan, on the share of customers owing proof to more than one rule maker. No published Canadian figure exists for that share, and the workbook carries a check that flags the line above fifteen per cent of revenue. The populations it is aimed at are thinner still: 400 larger property management brokerages and 3,100 smaller ones, both management estimates against no published count. And \$585,000 of authoring is budgeted for four rule sets rather than one. **The test is scheduled for the same quarter the seed is priced on, which is the worst possible time to discover the answer.**

RESPONSE FOR 4ORM

The worked example is already written and it is the right shape: a monthly broker trust check expressed as a job with a shape, a frequency, a due date, the two sides to compare, what it needs, who signs, a rule that whoever prepares it cannot be whoever signs it, and a retention period. Nothing in that has to be built, because the check shape already exists and frequency, due date and signer are fields rather than branches. **Run the test in 2027 on a duty chosen because it is awkward, not in 2028 on the one the expansion needs.** Before any of it is committed, require the rules lead and the senior engineer to produce jointly a line item accounting of what the build touches. If the seven things in the evidence store, the three shapes or the five screens need a new field, a new workflow state or a new binding rule, that is the signal, and it should force a hard decision well before Q3 2028 rather than a rushed one inside it. Switch the line off in every conversation until a second rule maker customer is signed, which the plan already recommends. Doing so improves the road: without Line E the lowest month moves from a shortfall of \$778,960 to \$624,143.

SOURCES ON THIS PAGE

Line E	\$698,381 in 2028 to \$13,766,769 in 2031, \$26,049,518 across the plan, on a share with no published Canadian figure. The workbook flags the line above fifteen per cent of revenue. 4orm Finance master pro forma rebuild of 16 August 2026, base case, recalculated in full, revenue lines tab.	MODEL
The populations	400 larger and 3,100 remaining property management brokerages, both management estimates against no published count. 4orm Finance master pro forma rebuild of 16 August 2026, base case, recalculated in full, population register tab.	MODEL
The worked example	The monthly broker trust check written entirely as content, and the four rule sets the \$585,000 buys. 4orm Finance Technical Architecture 03.2, 17 August 2026 · 4orm Finance master pro forma rebuild of 16 August 2026, base case, recalculated in full, development expense tab.	MODEL
The four rule sets	Alberta property management, plus three Ontario cadences that do not agree with each other: mortgage administrators file within 90 days of fiscal year end (FSRA), brokerages are reviewed monthly by the broker of record (RECO), and lawyers reconcile monthly within 25 days of month end (Law Society of Ontario).	PRIMARY
The milestone	Milestone 16, second rule maker configured with no second version of the software, Q3 2028. Milestone 18, seed closed on evidence, Q4 2028. 4orm Finance master pro forma rebuild of 16 August 2026, base case, recalculated in full, milestones tab.	MODEL

3

The held mechanism delays the two layers nobody else may build

SCORE 6 · L2 × I3

EVIDENCE · NINE FIELDS RELEASED, SIX ITEMS HELD PENDING A FILING

The record shape is agreed and published to the build partner tier: nine fields in three groups, answering which record and whose, which day and which rule, and who said so. A tenth and eleventh field sit in the third group and are not named. So does where the chain is anchored, how a replacing record is admitted into a store that permits no change, and how a record is serialised so that two exports a year apart land on identical bytes. An application covering the mechanism is in preparation with counsel and the items are released to a build partner on the day counsel confirms the filing. **The nine fields are enough to start the platform, the ingest and the core. They are not enough to build the two layers that are never delegated.**

RESPONSE FOR 4ORM

This is a sequencing risk rather than a technical one, and it has a date attached to somebody else's diary. The architecture note is due in Q3 2026 and gates the third instalment along with the intellectual property assignment. If counsel has not confirmed the filing by then, the work that can proceed is the work the partner does anyway, and the two layers that carry the whole claim wait. Three controls, none of them expensive. Put the counsel confirmation on the same status report as the two searches, weekly, so the date is tracked rather than assumed. Sequence the build so the partner's platform, ingest and core work fills the waiting period rather than idling against it. And write the disclosure step into the partner agreement now, under the existing non-disclosure agreement, so that confirmation triggers release automatically instead of starting a negotiation. **The assignment matters as much as the filing: everything built is assigned to 4orm Finance on creation, with a list of anything the partner already owned and a written waiver of the developer's personal claim to be credited, signed before any deposit moves.**

4

Re-export does not reproduce the same bytes

SCORE 4 · L2 × I2

EVIDENCE · THE ONE TEST THAT CANNOT BE ANSWERED WITH A QUESTIONNAIRE

Three things stop a record being changed in the months between when it is made and when it is asked for. There is no change path and no delete path anywhere in the code, blocked by the store itself rather than by a policy above it. Each record is locked to the one before it, so altering one shows up in plain sight. And the month can be exported again a year later and compared: the same bytes, or the store is not add only, whatever the paperwork says. Milestone 10 puts that last one on every build in Q2 2027 rather than in a security questionnaire. **It is also the only one of the three a customer's reviewer can verify without being shown the code.**

RESPONSE FOR 4ORM

Keep the test on every build and run it against production shaped feed data rather than synthetic fixtures, because the shape and volume of real bank and custodian files is where confidence usually outruns exposure. Then keep running it against increasingly old records rather than validating it once at launch: the value of the archive is that it existed on the day the event happened, and a test that only proves last month reproduces does not prove that. **Treat any non-identical output as stop-ship rather than as a known issue to patch, because the records this affects are the oldest ones, which are the ones an examiner opens with.** The commercial reason is in the comparables: nobody in the assessed field sells a record an outside reviewer can prove has gone untouched, and a record that carries its own proof is the one thing a competitor cannot answer by adding a rule badge to a framework list.

SOURCES ON THIS PAGE

The six boundaries	Feeds into the core, the rule store into the core, the core into the evidence store, the evidence store into the screens, the core into a filing, and the assistant into anything. 4orm Finance Technical Architecture 03.2, 17 August 2026, the six gates.	MODEL
The build split	Platform, ingest, the core, the screens and the assistant are built by the partner. The evidence store, the two dates, the rule store and the filing packs are ours, senior only. Everything built belongs to 4orm Finance the day it is made. 4orm Finance Technical Architecture 03.2, 17 August 2026.	MODEL
The duty the design answers	Section 15, a ledger showing the amount of funds belonging to each end user held at the end of each day. Justice Laws Canada · Bank of Canada, retail payments supervision	PRIMARY

5

The assistant writes into the record

SCORE 4 · L2 × I2

EVIDENCE · A DATED MODEL GOVERNANCE OBLIGATION ARRIVES AT THE BUYER'S BANK FIRST

The assistant sits behind the sixth boundary, and the boundary is written as a refusal: a search inside one customer with each answer tied to a record identifier may cross it, and a made up value entering the record may never. It finds and drafts for a person and it never writes proof. The reason to hold that line early is that somebody else's rule arrives before ours does. Model risk management guidance takes effect on 1 May 2027 for federally regulated financial institutions and covers models of this kind ([Office of the Superintendent of Financial Institutions](#)). It does not make a payment provider federally regulated. It does change what a bank counterparty asks a vendor sitting beside its payments ecosystem.

RESPONSE FOR 4ORM

Build the sign-off gate into the filing and examiner facing screens on day one rather than as a later feature. The screens are read only by design: the fourth boundary lets a read cross, logged as an event and scoped to one customer, and never a write of any kind, so the screens ask for a record and never make one. **The specific test for the design partner phase is that no workflow lets an assistant drafted answer reach a filing without a logged human approval in front of it.** Audit the tickets and the quality logs quarterly against it. The company boundary is the same one that answers the liability question in the regulatory domain, and it should be stated the same way every time: the duty stays with the firm holding the money, the customer's compliance officer signs every reconciliation and every filing, and this company never signs on a customer's behalf and never decides anything for them.

6

Ingest stops being configuration

SCORE 4 · L2 × I2

EVIDENCE · THE LAYER BUILT BY SOMEBODY ELSE, CARRYING A REVENUE LINE OF ITS OWN

Ingest takes files in from banks and other systems and checks them. It is built by the partner and reviewed before merge, no production data ever, and it is the layer most likely of the nine to accumulate work written for one customer. The first boundary is written as a refusal too: a checked record with a source and the time it arrived may cross, and a value patched up on the way in may not. If a field will not read, the record fails rather than defaults. **Connectors are also a named meter inside the retrieval and usage line, which carries \$17,035,003 in 2031 and is the least evidenced part of the build.** The same work is therefore modelled as revenue and exposed as cost, which is the reason to instrument it from the first integration.

RESPONSE FOR 4ORM

Track average engineering hours for each new feed as a standing number from the first design partner onward, and put it on the same report as days to first value. If it rises across 2027 rather than falling, the sentence the cost structure rests on is fraying at the layer built to hold configuration, and the answer is a mapping and detection framework rather than more configuration. Hold the refusal at the boundary while doing it: a field that will not read fails the record, and a defaulted value is never acceptable at any boundary for any reason, because a defaulted figure inside a daily position is precisely the kind of thing an outside reviewer is there to find. **Two revenue lines depend on this layer behaving, and the workbook already flags one of them: the meters check passes only because the line stays under a third of revenue, and the workbook says it is the assumption with the least support behind it.** Without the meters the lowest month on the road moves from a shortfall of \$778,960 to \$820,484, so they are additive rather than structural, which is the honest way to carry them.

SOURCES ON THIS PAGE

The held items	Nine fields released; the tenth and eleventh, the anchor, the admission of a replacing record and the serialisation held pending a filing with counsel. 4orm Finance Technical Architecture 03.2, 17 August 2026, the record shape.	MODEL
Gating milestones	Architecture note accepted and intellectual property assignment signed, both Q3 2026, both before the phase one deposit. Milestone 10, the automated re-export test, Q2 2027. 4orm Finance master pro forma rebuild of 16 August 2026, base case, recalculated in full, milestones tab.	MODEL
What stops a change	No change path and no delete path in the code, each record locked to the one before it, and proof by re-export. 4orm Finance Technical Architecture 03.2, 17 August 2026.	MODEL

What to watch, and how often

Every line asks to watch something run. A description of an architecture is not evidence of one, and the four items at the top of this list are all invisible on a demonstration, which is exactly why they are the ones that get traded away.

Monitoring cadence

- **The four things, at the architecture note and at every phase gate.** Demonstrated running rather than described: the same bytes on a second export, any day run again giving the same answer, a duty shipped with no code release, and a cross customer read returning nothing.
- **The seventh boundary, at every design review.** Ask for it. If a seventh place two parts touch exists and is not on the drawing, the drawing is wrong.
- **Counsel's confirmation of the filing, weekly, on the same report as the two searches.** The nine released fields start the platform, the ingest and the core. They do not start the two layers that are never delegated.
- **Every line item of the second rule maker estimate, well before Q3 2028.** A new field, a new workflow state or a new binding rule is the signal. Keep Line E switched off in conversation until a second rule maker customer is signed.
- **Assistant output paths, quarterly, from the design partner phase.** No assistant drafted answer reaches a filing without a logged human approval. Audit tickets and quality logs against that specific test.
- **Engineering hours for each new feed, from the first integration.** On the same report as days to first value. A rising trend across 2027 is the direct signal that configuration has become code.
- **Every build partner deliverable, before the next phase begins.** Against the internal architecture acceptance test rather than against a delivery date, and against the security review bought outside the base from a firm with no hand in the build.

SOURCES ON THIS PAGE

The boundaries	Boundary 1, feeds into the core: a value patched up on the way in must never cross, and a field that will not read fails the record. Boundary 4, the evidence store into the screens: a logged read, never a write. Boundary 6, the assistant: a search tied to a record identifier, never a made up value. 4orm Finance Technical Architecture 03.2, 17 August 2026.	MODEL
Model governance	Model risk management guidance takes effect 1 May 2027 for federally regulated financial institutions. Office of the Superintendent of Financial Institutions	PRIMARY
Line F	\$17,035,003 in 2031 across seven named meters including connectors. The workbook holds it under a third of revenue and states it is the assumption with the least support behind it. Without it the lowest month moves to a shortfall of \$820,484. 4orm Finance master pro forma rebuild of 16 August 2026, base case, recalculated in full · 4orm Finance Vision, 17 August 2026, section 13.	MODEL

The four things, priced

None of these four is dear at the start. All four are close to out of reach at the end, which is the whole case for spending the first release on the layers nobody can see. The test column is what a reviewer should ask to see running rather than described.

The thing	Cost now	What adding it later means	The test
A store that is only added to	Days	Every record written before the change has nothing behind it, so the trail is provable only from the day the change shipped. The oldest and most examined records are the ones left bare	The same bytes on a second export
Two dates on every record	One to two weeks	The rule versions that applied to past decisions were never kept, so they cannot be worked out from anything. The system can say what it believes now, which is not the question	Run any day again, get the same answer
Rules held as writing	Two to three weeks	Rules already in code have to be pulled out of it and every customer moved. In practice this gets put off, and the layer becomes a set of one off builds	Ship a duty with no code release
Customers apart at the key	Days, if designed in	Adding it later means re-keying live customer data. It is the most likely place for a breach to be let in, and it would be let in by the fix	A cross customer read returns nothing

Legend. Cost now is engineering time in the first release. 4orm Finance Technical Architecture 03.2, 17 August 2026.

Build the two hardest things first, because they are the two that cannot be added later.

Both are invisible in a demonstration and both are the whole reason a watched buyer would trust this company with the record they will be examined on. A new duty is a run of steps over the seven things already in the evidence store, never an eighth thing. The day somebody proposes an eighth table is the day the eleventh job stops being a setting, and the cost structure that follows from it stops being true.

SOURCES ON THIS PAGE

The design, and where it stops	4orm Finance Technical Architecture 03.2, 17 August 2026, read on a countersigned mutual non-disclosure agreement. 4orm Finance Technology Overview 03.1, 17 August 2026, open tier. Both trace to 4orm Finance Architecture Pack, 5 August 2026.	MODEL
Why it is the moat	Nobody in the assessed field sells a record an outside reviewer can prove has gone untouched. 4orm Finance Comparables and Benchmarks, 17 August 2026, section 5.	MODEL
Nothing in this budget	No blockchain, no token, no wallet, no custody, no settlement and no movement of money of any kind. The system reads balances and evidences them, and that boundary is what keeps this company outside the registration regime and outside the penalty ceilings that fall on its customers. 4orm Finance master pro forma rebuild of 16 August 2026, base case, recalculated in full, development expense tab.	MODEL
What a re-export reproduces	The day the record is about and the rule version that applied to it, both required by section 15, together with the reviewer independence requirement in section 17. Justice Laws Canada · reviewer independence · Bank of Canada safeguarding guideline	PRIMARY
Every figure on this document	4orm Finance master pro forma rebuild of 16 August 2026, base case, recalculated in full. 4orm Finance Master Source Register, 16 August 2026.	MODEL